

# 1 Methodology

Histogram calculation function is a simple for loop that takes the array and histogram array as inputs and fills the histogram array according to the number seen from the array that we generated.

## 1.1 Serial implementation

Pretty self-explanatory.

## 1.2 OpenMP implementation

### 1.2.1 Locks

Tried using locks to lock the interval before incrementing its count and then releasing after it's done. This gives a time that is actually more than the serial implementation. In fact, it performs around 50x worse than the simple serial implementation. This could be due to the fact that there are a very few intervals, which means a lot of contention for locks. Maybe this implementation could've worked well if we had a lot of intervals (in the thousands, perhaps).

### 1.2.2 Temporary local histogram array per thread

So we give a private array to each thread to increment its count on, so there is no contention between threads. When it's done, we add the temporary arrays together in a critical section. This should be fine since the code in the critical section is small.

# 2 Experimental Setup

Using an int array of length  $2^{32}$

So, size =  $2^2 \times 2^{30} \times 4B$

= 16 GB

the maximum size that I found the machine could allocate on the heap.

Loading the array with  $2^{32}$  values between 1 and 20 with seed=10

Using `clock_gettime()` function from the `time.h` library for more detailed timing and timing only the time spent in the histogram function.

## 3 Results

The correctness can be checked by looking at the output files and observing that the times are same for serial and parallel implementations regardless of the number of threads.

### 3.1 Serial Results

The serial implementation took time (in milliseconds) as seen in the table:

Interval size	1	2	3	4	5	Average
1	24710	24674	24792	24871	24710	24751
2	24447	24633	24557	24461	24447	24509
4	24489	24679	24603	24691	24489	24590
5	24636	24498	24537	24670	24636	24535
10	24482	24695	24671	24758	24482	24617

Average time = **24.6** seconds (regardless of interval size)

### 3.2 Parallel results

#### 3.2.1 Threads=2

Interval size	1	2	3	4	5	Average
1	15210	14338	14577	14430	14348	14580
2	14570	14417	14434	14373	14420	14442
4	14331	14375	14586	14376	14613	14456
5	14569	14395	14376	14526	14441	14461
10	14558	14440	14521	14465	14621	14521

Average time taken with 2 threads = **14.5** seconds (regardless of interval size)

Speedup = 1.7x

### 3.2.2 Threads=4

Interval size	1	2	3	4	5	Average
1	9259	9356	9119	9449	9253	9287
2	8970	9257	9193	9454	9348	9244
4	9262	9439	8894	9447	9246	9257
5	9235	9403	12154	9251	8997	9808
10	9458	9137	9368	9292	9194	9289

Average time taken with 4 threads = **9.3** seconds

Speedup = 2.6x

### 3.2.3 Threads=8

Interval size	1	2	3	4	5	Average
1	9333	9427	9146	10774	8657	9467
2	9632	10205	9066	8352	10011	9453
4	9196	10393	7594	8455	9365	9000
5	8702	7570	9072	10334	9620	9059
10	9655	9410	11163	10922	10762	10382

Average time taken with 8 threads = **9.2** seconds

Speedup = 2.7x

### 3.2.4 Threads=16 and onwards

The maximum number of threads that one node on the cluster can provide is 8. We can still go upto 16 threads by virtue of hyperthreading, but not beyond that. So we'll stick with 8 threads as maximum.

## 4 Observations

For an increasing number of threads, we observe a not-so-linear increase in speedup.




Figure 1: Average speedup obtained for different number of threads